

Frustrated?

Does Training take too long?

Does switching from CPU to
GPU cause you to have a mild
panic attack?

Are you lost in your own
code?!



Frustrated?

Does Training take too long?

Does switching from CPU to GPU cause you to have a mild panic attack?

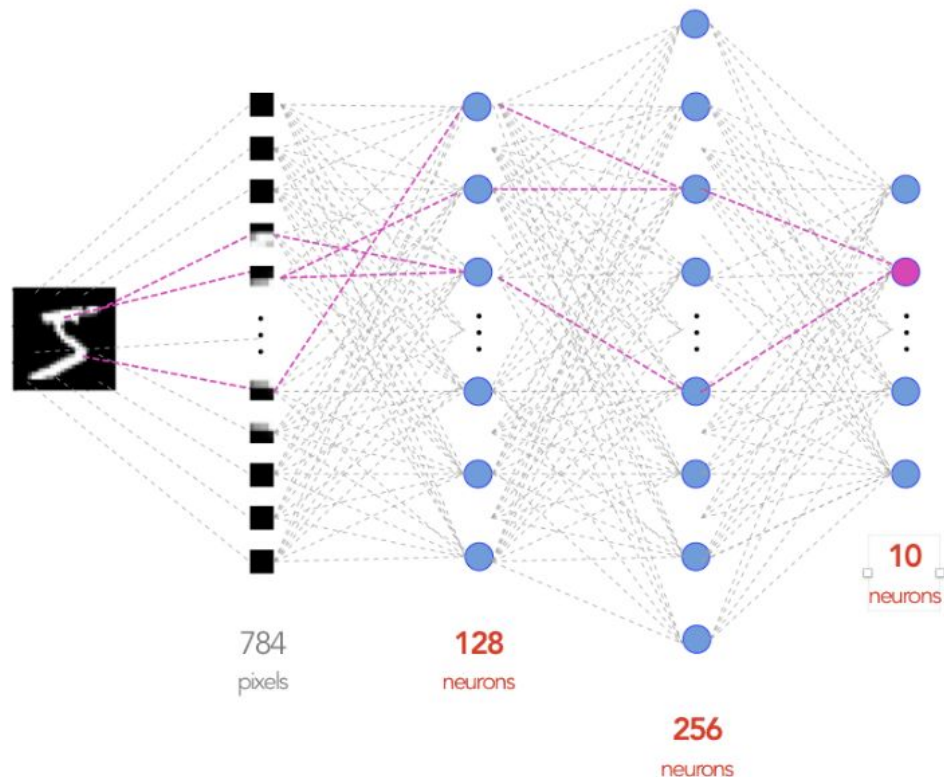
Are you lost in your own code?!



You need PyTorch Lightning!



You may ask yourself how would I implement a 3-layer fully connected neural network to perform classification over MNIST using PyTorch?



The Standard Project Approach

```
import torch
from torch import nn
from torch.utils.data import DataLoader, random_split
from torch.nn import functional as F
from torchvision.datasets import MNIST
from torchvision import datasets, transforms
import os

# -----
# MODEL
# -----
class MNISTClassifier(nn.Module):
    def __init__(self):
        super(MNISTClassifier, self).__init__()

        # mnist images are (1, 28, 28) (channels, width, height)
        self.layer_1 = torch.nn.Linear(28 * 28, 128)
        self.layer_2 = torch.nn.Linear(128, 256)
        self.layer_3 = torch.nn.Linear(256, 10)

    def forward(self, x):
        batch_size, channels, width, height = x.size()

        # (b, 1, 28, 28) -> (b, 1*28*28)
        x = x.view(batch_size, -1)

        # layer 1
        x = self.layer_1(x)
        x = torch.relu(x)

        # layer 2
        x = self.layer_2(x)
        x = torch.relu(x)

        # layer 3
        x = self.layer_3(x)

        # probability distribution over labels
        x = torch.log_softmax(x, dim=1)

        return x

# -----
# DATA
# -----
transform=transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.1307,), (0.3081,))])
mnist_train = MNIST(os.getcwd(), train=True, download=True, transform=transform)
mnist_test = MNIST(os.getcwd(), train=False, download=True, transform=transform)

# train (55,000 images), val split (5,000 images)
mnist_train, mnist_val = random_split(mnist_train, (55000, 5000))
mnist_test = MNIST(os.getcwd(), train=False, download=True)

# The dataloaders handle shuffling, batching, etc...
mnist_train = DataLoader(mnist_train, batch_size=64)
mnist_val = DataLoader(mnist_val, batch_size=64)
mnist_test = DataLoader(mnist_test, batch_size=64)

# -----
# OPTIMISER
# -----
pytorch_model = MNISTClassifier()
optimizer = torch.optim.Adam(pytorch_model.parameters(), lr=1e-3)

# -----
# LOSS
# -----
def cross_entropy_loss(logits, labels):
    return F.nll_loss(logits, labels)

# -----
# TRAINING LOOP
# -----
num_epochs = 1
for epoch in range(num_epochs):

    # TRAINING LOOP
    for train_batch in mnist_train:
        x, y = train_batch

        logits = pytorch_model(x)
        loss = cross_entropy_loss(logits, y)
        print('train loss: ', loss.item())

        loss.backward()

        optimizer.step()
        optimizer.zero_grad()

    # VALIDATION LOOP
    with torch.no_grad():
        val_loss = []
        for val_batch in mnist_val:
            x, y = val_batch
            logits = pytorch_model(x)
            val_loss.append(cross_entropy_loss(logits, y).item())

        val_loss = torch.mean(torch.tensor(val_loss))
        print('val_loss: ', val_loss.item())
```

The Model

Data

An Optimizer

The Loss

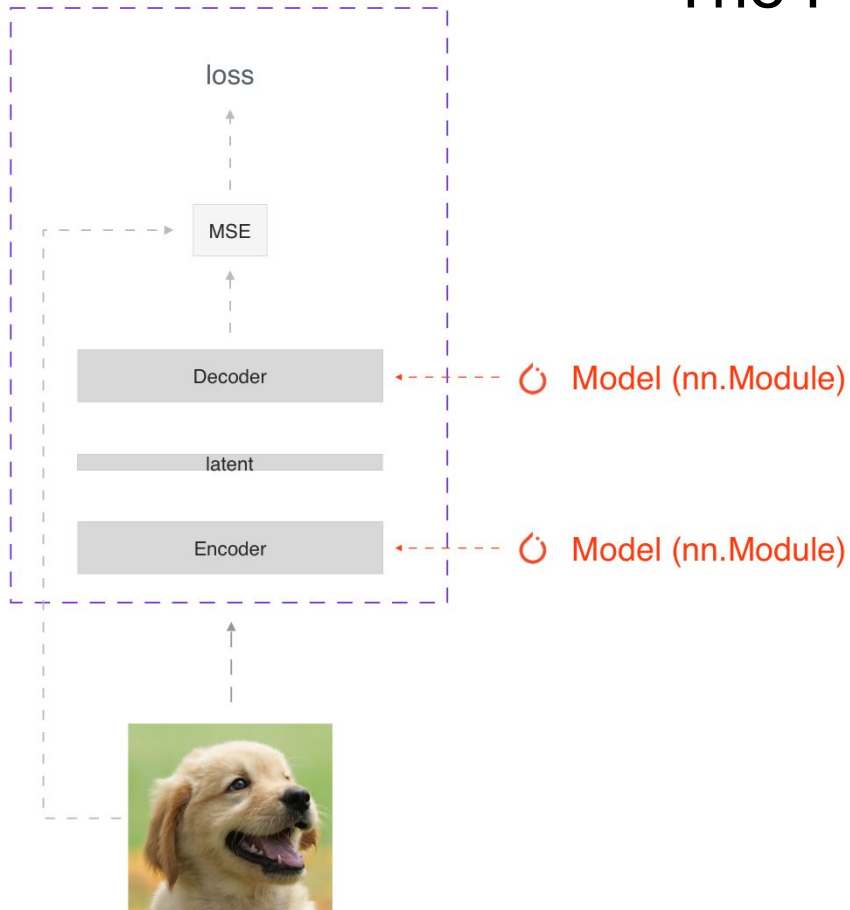
Training & Validation

Can we be more clear?

The PyTorch Lightning Approach

Define the system not the model.

- remove boilerplate code by decoupling research portion
- increase readability
- increase reproducibility
- scalable to any hardware!
- separates training from inference



PyTorch

```
from torch.utils.data import DataLoader, random_split
from torchvision.datasets import MNIST
import os
from torchvision import datasets, transforms

# -----
# TRANSFORMS
# -----
# prepare transforms standard to MNIST
transform=transforms.Compose([transforms.ToTensor(),
                             transforms.Normalize((0.1307,), (0.3081,))])

# -----
# TRAINING, VAL DATA
# -----
mnist_train = MNIST(os.getcwd(), train=True, download=True)

# train (55,000 images), val split (5,000 images)
mnist_train, mnist_val = random_split(mnist_train, [55000, 5000])

# -----
# TEST DATA
# -----
mnist_test = MNIST(os.getcwd(), train=False, download=True)

# -----
# DATALOADERS
# -----
# The dataloaders handle shuffling, batching, etc...
mnist_train = DataLoader(mnist_train, batch_size=64)
mnist_val = DataLoader(mnist_val, batch_size=64)
mnist_test = DataLoader(mnist_test, batch_size=64)
```

PyTorch Lightning

```
from torch.utils.data import DataLoader, random_split
from torchvision.datasets import MNIST
import os
from torchvision import datasets, transforms

class MNISTDataModule(pl.LightningDataModule)
    def prepare_data(self):
        # prepare transforms standard to MNIST
        MNIST(os.getcwd(), train=True, download=True)
        MNIST(os.getcwd(), train=False, download=True)

    def train_dataloader(self):
        transform=transforms.Compose([transforms.ToTensor(),
                                     transforms.Normalize((0.1307,), (0.3081,))])
        mnist_train = MNIST(os.getcwd(), train=True, download=False,
                           transform=transform)
        self.mnist_train, self.mnist_val = random_split(mnist_train, [55000, 5000])

        mnist_train = DataLoader(mnist_train, batch_size=64)
        return mnist_train

    def val_dataloader(self):
        mnist_val = DataLoader(self.mnist_val, batch_size=64)
        return mnist_val

    def test_dataloader(self):
        transform=transforms.Compose([transforms.ToTensor(),
                                     transforms.Normalize((0.1307,), (0.3081,))])
        mnist_test = MNIST(os.getcwd(), train=False, download=False,
                           transform=transform)
        mnist_test = DataLoader(mnist_test, batch_size=64)
        return mnist_test
```

prevents double manipulations to the data across all GPUs

Notice how we wrap the dataset train, validation, and test split in separate DataLoader member functions in a separate DataModule class.

PyTorch

```
[ ] import torch
    from torch import nn

    class MNISTClassifier(nn.Module):

        def __init__(self):
            super(MNISTClassifier, self).__init__()

            # mnist images are (1, 28, 28) (channels, width, height)
            self.layer_1 = torch.nn.Linear(28 * 28, 128)
            self.layer_2 = torch.nn.Linear(128, 256)
            self.layer_3 = torch.nn.Linear(256, 10)

        def forward(self, x):
            batch_size, channels, width, height = x.size()

            # (b, 1, 28, 28) -> (b, 1*28*28)
            x = x.view(batch_size, -1)

            # layer 1
            x = self.layer_1(x)
            x = torch.relu(x)

            # layer 2
            x = self.layer_2(x)
            x = torch.relu(x)

            # layer 3
            x = self.layer_3(x)

            # probability distribution over labels
            x = torch.log_softmax(x, dim=1)

            return x
```

PyTorch Lightning

```
[ ] import torch
    from torch import nn
    import pytorch_lightning as pl

    class LightningMNISTClassifier(pl.LightningModule):

        def __init__(self):
            super(LightningMNISTClassifier, self).__init__()

            # mnist images are (1, 28, 28) (channels, width, height)
            self.layer_1 = torch.nn.Linear(28 * 28, 128)
            self.layer_2 = torch.nn.Linear(128, 256)
            self.layer_3 = torch.nn.Linear(256, 10)

        def forward(self, x):
            batch_size, channels, width, height = x.size()

            # (b, 1, 28, 28) -> (b, 1*28*28)
            x = x.view(batch_size, -1)

            # layer 1
            x = self.layer_1(x)
            x = torch.relu(x)

            # layer 2
            x = self.layer_2(x)
            x = torch.relu(x)

            # layer 3
            x = self.layer_3(x)

            # probability distribution over labels
            x = torch.log_softmax(x, dim=1)

            return x
```

The Model: same code!

PyTorch

```
pytorch_model = MNISTClassifier()  
optimizer = torch.optim.Adam(pytorch_model.parameters(), lr=1e-3)
```

PyTorch Lightning

```
class LightningMNISTClassifier(pl.LightningModule):  
  
    def configure_optimizers(self):  
        optimizer = torch.optim.Adam(self.parameters(), lr=1e-3)  
        return optimizer
```

...and the Optimizer and Loss are the same too but are now member functions of the Model class!

PyTorch

```
from torch.nn import functional as F  
  
def cross_entropy_loss(logits, labels):  
    return F.nll_loss(logits, labels)
```

PyTorch Lightning

```
from torch.nn import functional as F  
  
class LightningMNISTClassifier(pl.LightningModule):  
  
    def cross_entropy_loss(self, logits, labels):  
        return F.nll_loss(logits, labels)
```


What about training?

In Pytorch if you...

- Go from CPU to using multiple GPUs
- Add gradient clipping
- Add early stopping
- Add checkpointing
- Use TPUs
- Use 16 bit precision



You will have to change your code and it can quickly start getting complex and difficult to understand.

STOP  **You need PyTorch Lightning!**

PyTorch

```

# -----
# TRAINING LOOP
# -----
num_epochs = 1
for epoch in range(num_epochs):

    # TRAINING LOOP
    for train_batch in mnist_train:
        x, y = train_batch

        logits = pytorch_model(x)
        loss = cross_entropy_loss(logits, y)
        print('train_loss: ', loss.item())

        loss.backward()
        optimizer.step()
        optimizer.zero_grad()

    # VALIDATION LOOP
    with torch.no_grad():
        val_loss = []
        for val_batch in mnist_val:
            x, y = val_batch
            logits = pytorch_model(x)
            val_loss = cross_entropy_loss(logits, y).item()
            val_loss.append(val_loss)

        val_loss = torch.mean(torch.tensor(val_loss))
        print('val_loss:', val_loss.item())

```

PyTorch Lightning

```

class LightningMNISTClassifier(pl.LightningModule):

    def training_step(self, train_batch, batch_idx):
        x, y = train_batch
        logits = self.forward(x)
        loss = self.cross_entropy_loss(logits, y)
        self.log('train_loss', loss)
        return loss

    def validation_step(self, val_batch, batch_idx):
        x, y = val_batch
        logits = self.forward(x)
        loss = self.cross_entropy_loss(logits, y)
        self.log('val_loss', loss)

```

(automatically reduced across epochs)

Now that's **readable!**

To run training use **Trainer**:

PyTorch

```
# -----  
# TRAINING LOOP  
# -----  
num_epochs = 1  
for epoch in range(num_epochs):  
  
    # TRAINING LOOP  
    for train_batch in mnist_train:  
        x, y = train_batch  
  
        logits = pytorch_model(x)  
        loss = cross_entropy_loss(logits, y)  
        print('train loss: ', loss.item())  
  
        loss.backward()  
  
        optimizer.step()  
        optimizer.zero_grad()  
  
    # VALIDATION LOOP  
    with torch.no_grad():  
        val_loss = []  
        for val_batch in mnist_val:  
            x, y = val_batch  
            logits = pytorch_model(x)  
            val_loss.append(cross_entropy_loss(logits, y).item())  
  
        val_loss = torch.mean(torch.tensor(val_loss))  
        print('val_loss: ', val_loss.item())
```

PyTorch Lightning

```
# train loop + val loop + test loop  
trainer = pl.Trainer()  
trainer.fit(LightningMNISTClassifier())
```

Look, only two lines, wow!

Notice that the code

- doesn't loop over epochs or datasets
- doesn't set the model to evaluation or training
- doesn't enable or disable gradients!

The **Trainer** also has flags such as

- auto_scale_batch_size
- auto_lr_find
- fast_dev_run

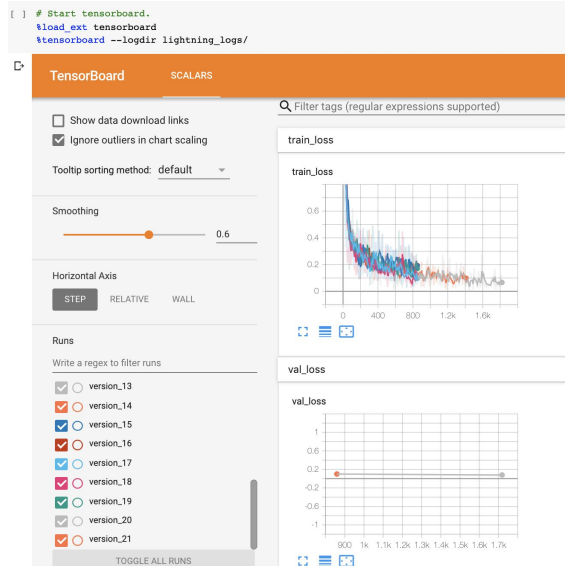
...to name a few...

But wait there's more...

- Colorful progress bars
- weights summary
- tensorboard logs
- training on CPUs, GPUs, or TPUs without changing your code
- 16 bit precision training to speed things up
- a profiler
- extensibility with hooks



	Name	Type	Params
0	layer_1	Linear	100 K
1	layer_2	Linear	33 K
2	layer_3	Linear	2 K



And Callbacks

A callback is non-essential code that is not related to the research portion:

```
import pytorch_lightning as pl

class MyPrintingCallback(pl.Callback):

    def on_init_start(self, trainer):
        print('Starting to init trainer!')

    def on_init_end(self, trainer):
        print('trainer is init now')

    def on_train_end(self, trainer, pl_module):
        print('do something when training ends')
```

```
trainer = pl.Trainer(num_tpu_cores=8, callbacks=[MyPrintingCallback()])
trainer.fit(model)
```




And just like that you cleaned your code into:

1. DataLoader code (DataModule)
2. Research code (LightningModule)
3. Engineering code (Trainer)
4. Non-research code (Callbacks)



****Terms and Conditions May Apply:**

PyTorch Lightning is good for rudimentary off-the-shelf things like boilerplate code for training but may hide details specific to more complex models.. PyTorch Lightning is not recommended for prototyping more complex, application-specific cases.